

# CB2-101: R for Bioinformatics

Malay ([malay@uab.edu](mailto:malay@uab.edu))

November 14, 2016

## Contents

<b>1</b>	<b>Some useful resources</b>	<b>1</b>
1.1	BAM or SAM format . . . . .	2
1.2	VCF . . . . .	2
<b>2</b>	<b>Where to start ?</b>	<b>2</b>
<b>3</b>	<b>A simple example</b>	<b>2</b>
<b>4</b>	<b>Some basic objects in BioConductor</b>	<b>3</b>
4.1	IRanges . . . . .	3
4.1.1	Simple operations on IRanges . . . . .	3
4.1.2	Getting the flanking region . . . . .	4
4.1.3	Set operations on ranges . . . . .	4
4.2	Run length Encoding (RLE) . . . . .	5
4.3	GenomicRanges . . . . .	5
4.3.1	GRanges . . . . .	5
4.3.2	GRangesList . . . . .	6
4.3.3	GappedAlignments . . . . .	7
4.4	BSgenome . . . . .	7
<b>5</b>	<b>Annotation database</b>	<b>7</b>
5.1	What is the mutation frequency of P53 gene in normal human population . . . . .	10

## 1 Some useful resources

R was a popular tool for analysis of microarray data. Now it is mostly used in Bioinformatics for analysis of next-gen sequence data. It is not very popular as a general purpose Bioinformatics tool. There are a bunch of special packages distributed under the name “BioConductor” (<http://www.bioconductor.org/>) that are related to biological data analysis using R.

- Bioinformatics using R
  1. A little Book of R for Bioinformatics (<http://a-little-book-of-r-for-bioinformatics.readthedocs.org/en/latest/>).
- Learning BioConductor
  1. BioConductor help section contains exhaustive lists of conferences (<http://www.bioconductor.org/help/course-materials/>). The course materials of these conferences are very good.
  2. A nice intermediate level guide to R and BioConductor: <http://www.bioconductor.org/help/course-materials/2013/SeattleMay2013/IntermediateSequenceAnalysis2013.pdf>.
  3. A somewhat scattered introduction to NGS data analysis using R and BioConductor: <http://manuals.bioinformatics.ucr.edu/home/ht-seq>

You can see there is sharp drop of quality below score 20. This is why Phred 20 is a good cutoff score. This actually  $(20 + 33) = 53$  which 5 in ascii.

## 1.1 BAM or SAM format

The FASTQ files are aligned against a reference genome using a software like BWA (<http://bio-bwa.sourceforge.net/>). The resulting alignment format is a BAM or SAM files. BAM files are binary, SAM files are plain text. The software for interconversion and analysis of these files are mainly `samtools` (<http://www.htslib.org/>). A small example BAM files comes along with `Rsamtools` package. Sam file format specification can be found here <http://samtools.github.io/hts-specs/SAMv1.pdf>.

## 1.2 VCF

Once the alignment BAM files have been generated, a variant caller like GATK (<https://www.broadinstitute.org/gatk/>) is used to find the variants in the file. The resulting file is called VCF. The specification can be found here (<http://samtools.github.io/hts-specs/VCFv4.2.pdf>). A sample VCF line is given below:

```
chr1    873762    .        T        G    [CLIPPED] GT:AD:DP:GQ:PL    0/1:173,141:282:99:255,0,255
chr1    877664    rs3828047 A        G    [CLIPPED] GT:AD:DP:GQ:PL    1/1:0,105:94:99:255,255,0
chr1    899282    rs28548431 C        T    [CLIPPED] GT:AD:DP:GQ:PL    0/1:1,3:4:25.92:103,0,26
```

Once the variant is called they are annotated using variant annotation tools like SnpEff (<http://snpeff.sourceforge.net/>) or Annovar (<http://www.openbioinformatics.org/annovar/>) or `VariantAnnotation` package.

```
-> # Installing BioConductor packages
```

All bioconductor packages are installed using the following commands:

```
source("http://bioconductor.org/biocLite.R")
biocLite("packagename")
```

Where, `packagename` is the name of your BioConductor package.

## 2 Where to start ?

BioConductor is a jumbled mess of hundreds of packages. And a problem for the beginners is to know where to start and which packages to use. I suggest you start with the workflows page of BC (<http://bioconductor.org/help/workflows/>). Look at the examples and find out what packages are used and then go and dig for more information about those packages.

## 3 A simple example

Lets start with a simple example. Remember, we calculated the average protein length of *E. coli* in our Linux problem set. Let's solve this using BC. The package that we need is `Biostrings`. Let's install the package.

```
source("http://bioconductor.org/biocLite.R")
biocLite("Biostrings")
```

Once the package is installed. We have to now load it.

```
suppressPackageStartupMessages( library("Biostrings") )
```

You can see an overview of what `Biostrings` package has to offer.

```
browseVignettes("Biostrings")
```

You can now get a quick overview by clicking on “Biostrings quick overview” PDF. By looking at the quick overview, we find that there is a function in `Biostrings` that can read the sequence: `readAAStringSet()`.

```
# Just to get the long line to wrap correctly
url <- paste("ftp://ftp.ncbi.nlm.nih.gov/genomes/archive/old_refseq/Bacteria/Escherichia_coli_K_12_subs
faa <- readAAStringSet(url)
```

We can get the average length now:

```
av.length <- sum( width(faa) )/length(faa)
av.length
```

```
## [1] 316.8587
```

## 4 Some basic objects in BioConductor

### 4.1 IRanges

IRanges represents ordered indices.

```
library(IRanges)
r <- IRanges(start=c(1,3,12,10),end=c(4,5,25,19))
r
```

```
## IRanges of length 4
##      start end width
## [1]     1  4     4
## [2]     3  5     3
## [3]    12 25    14
## [4]    10 19    10
```

#### 4.1.1 Simple operations on IRanges

```
length(r)
```

```
## [1] 4
```

```
start(r)
```

```
## [1] 1 3 12 10
```

```
end(r)
```

```
## [1] 4 5 25 19
```

```
width(r)
```

```
## [1] 4 3 14 10
```

```
r[1:2]
```

```
## IRanges of length 2
##      start end width
## [1]     1  4     4
## [2]     3  5     3
```

```
range(r)
```

```
## IRanges of length 1
##      start end width
```

```
## [1]      1 25    25
```

```
reduce(r)
```

```
## IRanges of length 2
```

```
##      start end width
```

```
## [1]      1  5     5
```

```
## [2]     10 25    16
```

```
disjoin(r)
```

```
## IRanges of length 6
```

```
##      start end width
```

```
## [1]      1  2     2
```

```
## [2]      3  4     2
```

```
## [3]      5  5     1
```

```
## [4]     10 11     2
```

```
## [5]     12 19     8
```

```
## [6]     20 25     6
```

```
coverage(r)
```

```
## integer-Rle of length 25 with 7 runs
```

```
##   Lengths: 2 2 1 4 2 8 6
```

```
##   Values : 1 2 1 0 1 2 1
```

#### 4.1.2 Getting the flanking region

```
flank(r, 1, both=T, start=T)
```

```
## IRanges of length 4
```

```
##      start end width
```

```
## [1]      0  1     2
```

```
## [2]      2  3     2
```

```
## [3]     11 12     2
```

```
## [4]      9 10     2
```

#### 4.1.3 Set operations on ranges

```
r2 <- IRanges(start=c(7,8,14),end=c(11,16,18))
```

```
union(r,r2)
```

```
## IRanges of length 2
```

```
##      start end width
```

```
## [1]      1  5     5
```

```
## [2]      7 25    19
```

```
intersect(r,r2)
```

```
## IRanges of length 1
```

```
##      start end width
```

```
## [1]     10 18     9
```

```
setdiff(r,r2)
```

```
## IRanges of length 2
```

```
##      start end width
## [1]      1  5      5
## [2]     19 25      7
```

## 4.2 Run length Encoding (RLE)

```
x<- Rle(c(1,1,2,2,2))
length(x)
```

```
## [1] 5
```

```
start(x)
```

```
## [1] 1 3
```

```
end(x)
```

```
## [1] 2 5
```

```
width(x)
```

```
## [1] 2 3
```

```
nrun(x)
```

```
## [1] 2
```

```
runLength(x)
```

```
## [1] 2 3
```

## 4.3 GenomicRanges

There are 3 classes in this package: `GRanges`, `GRangeList`, `GappedAlignments`.

### 4.3.1 GRanges

```
library(GenomicRanges)
```

```
## Loading required package: GenomeInfoDb
```

```
gr <- GRanges(seqnames= Rle(c("chr1", "chr2"), c(2, 3)),
              ranges = IRanges(1:5, end= 6:10),
              strand = Rle(strand(c("-", "+", "+", "-", "+"))),
              score=1:5, GC=seq(1, 0, length=5))
gr
```

```
## GRanges object with 5 ranges and 2 metadata columns:
##      seqnames   ranges strand |      score      GC
##      <Rle> <IRanges> <Rle> | <integer> <numeric>
## [1]   chr1   [1, 6]    - |         1         1
## [2]   chr1   [2, 7]    + |         2        0.75
## [3]   chr2   [3, 8]    + |         3         0.5
## [4]   chr2   [4, 9]    - |         4        0.25
## [5]   chr2   [5, 10]   + |         5          0
## -----
```

```
## seqinfo: 2 sequences from an unspecified genome; no seqlengths
```

#### 4.3.1.1 Access elements of GRanges

```
length(gr)
```

```
## [1] 5
```

```
seqnames(gr)
```

```
## factor-Rle of length 5 with 2 runs
##   Lengths:    2    3
##   Values  : chr1 chr2
## Levels(2): chr1 chr2
```

```
start(gr)
```

```
## [1] 1 2 3 4 5
```

```
end(gr)
```

```
## [1] 6 7 8 9 10
```

```
ranges(gr)
```

```
## IRanges of length 5
##      start end width
## [1]     1  6     6
## [2]     2  7     6
## [3]     3  8     6
## [4]     4  9     6
## [5]     5 10     6
```

```
strand(gr)
```

```
## factor-Rle of length 5 with 4 runs
##   Lengths: 1 2 1 1
##   Values  : - + - +
## Levels(3): + - *
```

All other fields besides `seqnames`, `range` and `strands` need to be accessed by `elementMetadata` function.

```
elementMetadata(gr)
```

```
## DataFrame with 5 rows and 2 columns
##      score      GC
##   <integer> <numeric>
## 1         1      1.00
## 2         2      0.75
## 3         3      0.50
## 4         4      0.25
## 5         5      0.00
```

#### 4.3.2 GRangesList

It's a list of `GRanges` objects.

```
GRangesList (gr, gr)
```

```
## GRangesList object of length 2:
## [[1]]
## GRanges object with 5 ranges and 2 metadata columns:
##      seqnames      ranges strand |      score      GC
##      <Rle> <IRanges> <Rle> | <integer> <numeric>
## [1]   chr1    [1,  6]     - |         1         1
## [2]   chr1    [2,  7]     + |         2        0.75
## [3]   chr2    [3,  8]     + |         3         0.5
## [4]   chr2    [4,  9]     - |         4        0.25
## [5]   chr2    [5, 10]     + |         5          0
##
## [[2]]
## GRanges object with 5 ranges and 2 metadata columns:
##      seqnames      ranges strand | score  GC
## [1]   chr1 [1,  6]     - |    1    1
## [2]   chr1 [2,  7]     + |    2 0.75
## [3]   chr2 [3,  8]     + |    3 0.5
## [4]   chr2 [4,  9]     - |    4 0.25
## [5]   chr2 [5, 10]     + |    5    0
##
## -----
## seqinfo: 2 sequences from an unspecified genome; no seqlengths
```

### 4.3.3 GappedAlignments

Used for parsing BAM files.

## 4.4 BSgenome

**BSgenome** is the actual genome sequences distributed in a R package. This packages can be pretty big. For human this file is about 1.7G in size. For this course, we will not use it anymore.

## 5 Annotation database

There are two types of annotation databases in BC. Organism-specific gene level databases are names as `org.XX.XXX.db`. For e.g., `org.Hs.eg.db`. This is human entrez gene database. There are also metapackages (not for all organisms) that pull data from may different sources. `Homo.sapiens` is one such databases. Let's use this database.

```
suppressPackageStartupMessages(library("Homo.sapiens"))
columns(Homo.sapiens)
```

```
## [1] "ACCNUM"      "ALIAS"      "CDSCHROM"   "CDSEND"
## [5] "CDSID"      "CDSNAME"   "CDSSTART"   "CDSSTRAND"
## [9] "DEFINITION" "ENSEMBL"    "ENSEMBLPROT" "ENSEMBLTRANS"
## [13] "ENTREZID"   "ENZYME"     "EVIDENCE"    "EVIDENCEALL"
## [17] "EXONCHROM"  "EXONEND"    "EXONID"      "EXONNAME"
## [21] "EXONRANK"   "EXONSTART"  "EXONSTRAND"  "GENEID"
## [25] "GENENAME"   "GO"         "GOALL"       "GOID"
## [29] "IPI"        "MAP"        "OMIM"        "ONTOLOGY"
## [33] "ONTOLOGYALL" "PATH"       "PFAM"        "PMID"
## [37] "PROSITE"    "REFSEQ"     "SYMBOL"      "TERM"
```

```
## [41] "TXCHROM"      "TXEND"      "TXID"      "TXNAME"
## [45] "TXSTART"     "TXSTRAND"   "TXTYPE"    "UCSCKG"
## [49] "UNIGENE"     "UNIPROT"
```

Only some of these columns can be use to retrieve data. To find what columns can be used

```
keytypes(Homo.sapiens)
```

```
## [1] "ACCNUM"      "ALIAS"      "CDSID"      "CDSNAME"
## [5] "DEFINITION" "ENSEMBL"    "ENSEMBLPROT" "ENSEMBLTRANS"
## [9] "ENTREZID"    "ENZYME"     "EVIDENCE"    "EVIDENCEALL"
## [13] "EXONID"      "EXONNAME"   "GENEID"      "GENENAME"
## [17] "GO"          "GOALL"      "GOID"        "IPI"
## [21] "MAP"         "OMIM"       "ONTOLOGY"    "ONTOLOGYALL"
## [25] "PATH"        "PFAM"       "PMID"        "PROSITE"
## [29] "REFSEQ"      "SYMBOL"     "TERM"        "TXID"
## [33] "TXNAME"      "UCSCKG"     "UNIGENE"     "UNIPROT"
```

To extract data we need the “keys” corresponding to a “keytype”. For example the SYMBOL keytypes stores the gene name and surprisingly GENENAME actually contains a description of gene. We can show the partial list of these genes.

```
genenames<-(keys(Homo.sapiens,keytype="SYMBOL"))
```

There are altogether 60083 genes in this database. We can now use genenames as keys to get the genes and their longer name for the database.

```
gene.list <-select(Homo.sapiens,keys=genenames,columns=c("SYMBOL","GENENAME"),keytype="SYMBOL")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
head(gene.list)
```

```
##      SYMBOL      GENENAME
## 1    A1BG      alpha-1-B glycoprotein
## 2    A2M      alpha-2-macroglobulin
## 3  A2MP1  alpha-2-macroglobulin pseudogene 1
## 4   NAT1      N-acetyltransferase 1
## 5   NAT2      N-acetyltransferase 2
## 6   NATP  N-acetyltransferase pseudogene
```

Let’s do something interesting. Let plot the number of genes per chromosomes.

```
# Old bioconductor gene.df <-
# select(Homo.sapiens,keys=genenames,columns=c('CHROM','CHRLOC','CHRLOCEND'),keytype='SYMBOL')
gene.df <- select(Homo.sapiens, keys = genenames, columns = c("TXCHROM"), keytype = "SYMBOL")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
head(gene.df)
```

```
##      SYMBOL TXCHROM
## 1    A1BG    chr19
## 2    A2M    chr12
## 3  A2MP1    chr12
## 4   NAT1    chr8
## 5   NAT2    chr8
## 6   NATP    <NA>
```




```
# Let's extract the SYMBOL and CHR is a separate dataframe.
gene.uniq <- data.frame(symbol = gene.df$SYMBOL, chr = gene.df$TXCHROM)
```

```
# Let's remove the duplicated lines.
gene.uniq <- gene.uniq[order(gene.uniq$symbol), ]
gene.uniq <- gene.uniq[!duplicated(gene.uniq), ]
head(gene.uniq)
```

```
##      symbol  chr
## 38980 1060P11.3 <NA>
## 23563   3.8-1.2 <NA>
## 23564   3.8-1.3 <NA>
## 23565   3.8-1.4 <NA>
## 23566   3.8-1.5 <NA>
## 24362   5-HT3C2 <NA>
```

```
gene.freq <- table(gene.uniq$chr)
barplot(table(gene.uniq$chr), las = 2)
```



```
## OK
```

```
head(chr.info)
```

```
##   chrom   length
## 1  chr1 249250621
## 2  chr2 243199373
## 3  chr3 198022430
## 4  chr4 191154276
## 5  chr5 180915260
## 6  chr6 171115067
```

```
# Convert our frequency table into data frame
```

```
gene.freq <- data.frame(gene.freq)
names(gene.freq) <- c("chr", "freq")
```

```
# We need to convert the names of the chr column
```

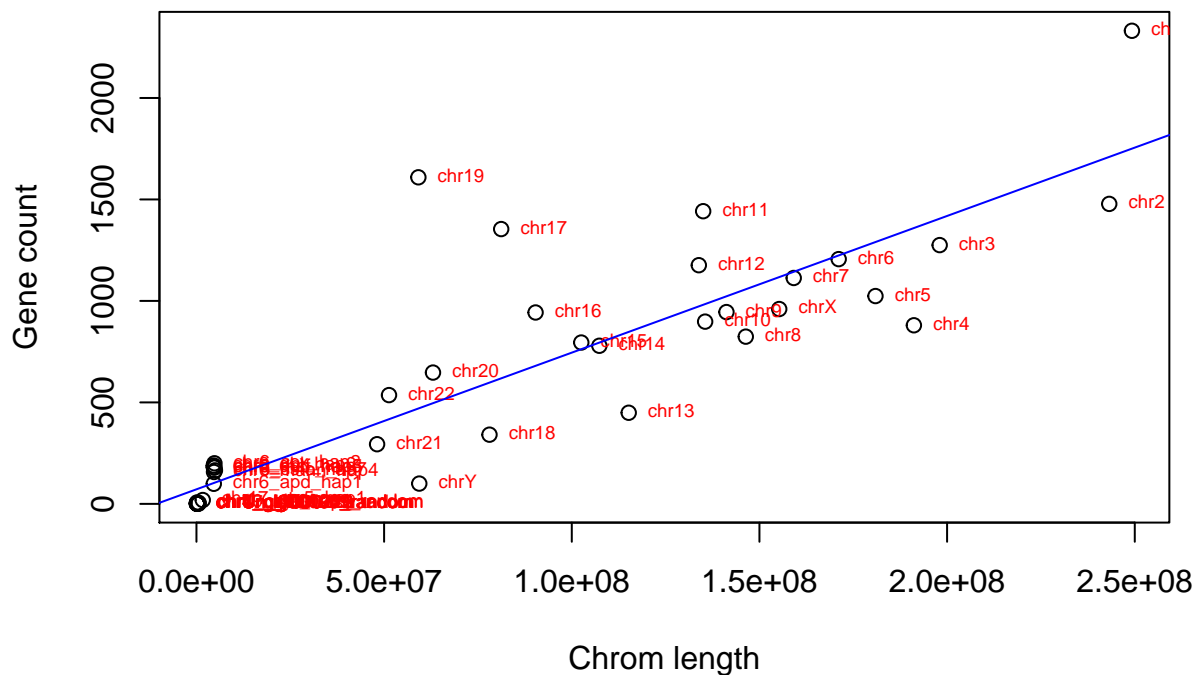
```
#gene.freq$chr <- paste('chr', gene.freq$chr, sep="")
```

```
merged.data <- merge(gene.freq, chr.info, by.x="chr", by.y="chrom")
```

```
plot(merged.data$length, merged.data$freq, xlab="Chrom length", ylab="Gene count")
```

```
text(merged.data$length, merged.data$freq, merged.data$chr, cex=0.6, pos=4, col="red")
```

```
abline(lm(merged.data$freq~merged.data$length), col="blue")
```



## 5.1 What is the mutation frequency of P53 gene in normal human population

For this problem we first have to find the location of the P53 gene in human annotation database.

```
library(Homo.sapiens)
loc <- select(Homo.sapiens,keys="TP53",columns=c("SYMBOL","CHR","CHRLOC","CHRLOCEND"),keytype="SYMBOL")
```

## 'select()' returned 1:1 mapping between keys and columns

We see that TP53 gene is on chromosome 17 in location 7571720:7590868. We will download this portion of the variation from 1000 genome data using **tabix**. Install **tabix** on your system.

Once **tabix** is installed. We can download this portion of the file using the following command.

```
tabix -fh ftp://ftp-trace.ncbi.nih.gov/1000genomes/ftp/release/20110521/\
ALL.chr17.phase1_release_v3.20101123.snps_indels_sv.s.genotypes.vcf.gz \
17:7571720-7590868 >p53.vcf
```

Let's read the VCF file in R.

```
library(VariantAnnotation)
vcf <- readVcf("p53.vcf","hg19")
```

We will now locate variant using the txdb package.

```
library("TxDb.Hsapiens.UCSC.hg19.knownGene")
txdb <- TxDb.Hsapiens.UCSC.hg19.knownGene
txdb <- renameSeqlevels(txdb, gsub("chr","",seqlevels(txdb)))
txdb <- keepSeqlevels(txdb,"17")
all <- locateVariants(vcf,txdb, AllVariants())
table(mcols(all)$LOCATION)
```

Looks like there are 162 variants in the coding regions in the 1000K sample. I will leave it to you to investigate this further.

---